The Institutional Semantic Observatory

Engineering Information Environments for Explainable Institutional AI

Edward Brash

Christopher Newport University
School of Engineering and Computing

July 13, 2026

"Information is the resolution of uncertainty."

— Claude Shannon

Universities Have an Information Problem

Modern universities accumulate enormous quantities of institutional knowledge.

- Policies
- Faculty records
- Budgets
- Strategic plans
- Assessment reports
- Accreditation documents
- Meeting minutes
- Enrollment reports
- Facilities inventories
- Historical archives

A medium-sized university may possess

$$15{,}000+$$

documents spanning decades.

A Thought Experiment

Suppose the Provost asks:

"Should CNU create a Mechanical Engineering major?"

An experienced dean would begin by gathering evidence.

- Faculty expertise
- Laboratory space
- Equipment
- ABET implications
- Budget
- Enrollment projections
- Strategic priorities

Observation

Institutional reasoning begins with evidence gathering. Only then does reasoning occur.

The Human Bottleneck

Imagine giving an experienced dean access to

15,000 institutional documents.

The dean faces two fundamental limitations.

1. **Scale**

   Reading the complete corpus could require months or years.

2. **Recall**

   Even with unlimited time,
   which ten documents actually matter?

Key Insight

The bottleneck is not human intelligence.
The bottleneck is locating the appropriate evidence.

Traditional Retrieval-Augmented Generation (RAG)

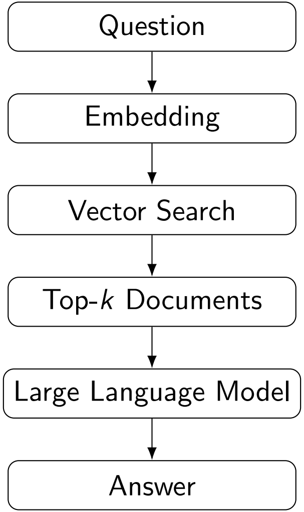

Where Does Intelligence Live?

Conventional AI View

- Intelligence resides primarily in the model.
- Better models should produce better answers.
- Guardrails constrain model behavior.
- Trust is placed mainly in the reasoning engine.

ISO View

- Intelligence emerges from the interaction between the model and its information environment.
- Better information environments produce better evidence.
- Guardrails shape the conditions under which reasoning occurs.
- Trust is placed in the complete process.

Central Thesis

Reliable institutional reasoning is an emergent property of a well-engineered information environment.

Two Philosophies of AI System Design

Conventional Approach

Build a more capable model

↓

Prevent it from doing bad things

↓

Produce a better answer

ISO Approach

Engineer the information environment

↓

Retrieve reliable evidence

↓

Reason over that evidence

↓

Support a defensible decision

One constrains the actor. The other engineers the environment.

Rethinking Guardrails

The word *guardrail* is often used as though it describes a single technical mechanism.

In practice, guardrails can be implemented at many layers.

- Corpus inclusion and exclusion rules
- Parser validation
- Provenance requirements
- Duplicate detection
- Chunking policy
- Retrieval constraints
- Evidence coverage checks
- Behavioral instructions
- Human approval requirements

Design Principle

Guardrails should be implemented at the lowest layer where they can be made explicit, observable, and auditable.

The most effective guardrail is a well-engineered information environment.

Why?

A model is less likely to produce an unjustified conclusion when it receives relevant, diverse, well-provenanced evidence and is explicitly told when that evidence is incomplete.

From Information Engineering to Semantic Engineering

Information Engineering

- Parse files
- Normalize text
- Preserve metadata
- Record provenance
- Store structured objects
- Build indexes

Semantic Engineering

- Measure corpus balance
- Detect dominant documents
- Identify semantic gaps
- Monitor retrieval quality
- Estimate evidence coverage
- Evaluate reasoning support

Information engineering organizes data.
Semantic engineering organizes the conditions for reasoning.

Knowledge Objects Store Facts

Knowledge Objects

A knowledge object records what is known about a source without embedding a conclusion into the source itself.

Typical fields include

- source path
- title
- parser
- timestamps
- document type
- normalized text
- semantic chunks
- provenance

Services Derive Meaning

Services

Services operate over knowledge objects to derive context-dependent meaning.

Examples include

- retrieval
- reranking
- evidence synthesis
- corpus diagnostics
- scenario analysis
- decision brief generation
- conflict detection
- confidence assessment

The ISO Pipeline

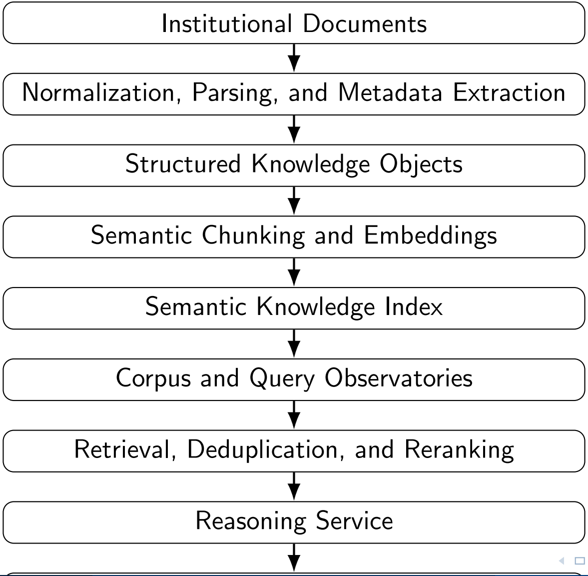

The Semantic Ecosystem

A corpus is not merely a collection of files.

It is an information environment with structure.

Healthy Properties

- Diversity
- Balance
- Provenance
- Connectivity
- Resilience
- Observability
- Evolution over time

Pathologies

- Dominant documents
- Duplicates
- Corrupted files
- Stale knowledge
- Semantic islands
- Missing regions
- Parser failures

Why an Observatory?

In every mature scientific system, instrumentation is used to monitor the health of the environment in which inference occurs.

Experimental Physics

Detector monitoring
Calibration
Efficiency
Systematics

Ecology

Population density
Diversity
Fragmentation
Resilience

ISO

Corpus balance
Parser health
Retrieval quality
Evidence support

The observatory is not an auxiliary dashboard.
It is part of the reasoning infrastructure.

The Corpus Observatory

The current observatory measures global properties of the corpus.

- Document count
- Total semantic chunks
- Chunks per document
- File-type distribution
- Parser distribution
- Folder distribution
- Dominance metrics
- Gini coefficient
- Largest documents
- Outlier detection

Semantic Observatory Screenshot

Insert corpus-health image here

What Does the Gini Coefficient Measure?

The Gini coefficient measures how unevenly semantic chunks are distributed across documents.

Low Gini

- Information is distributed relatively evenly.
- No small group of documents dominates the corpus.
- Retrieval is less likely to be overwhelmed by a few sources.

High Gini

- A small number of documents contain a large fraction of all chunks.
- Large or duplicated files may dominate retrieval.
- Corpus health may require investigation.

Important

The Gini coefficient is not a score to maximize or minimize.
It is a diagnostic indicator of corpus structure.

A Real Corpus Pathology

One spreadsheet initially generated approximately

50,000

semantic chunks.

The spreadsheet appeared modest when viewed normally.
However, several corrupted rows extended across thousands of columns and repeated the same values many times.

What the Observatory Revealed

- Extreme document dominance
- Inflated chunk counts
- Distorted corpus statistics
- Wasted embedding and indexing resources

From Corpus Health to Query Health

The Corpus Observatory asks:

Is the overall information environment healthy?

The next phase asks a more local question:

How well supported is this particular reasoning task?

Query Observatory

A Query Observatory evaluates the state of the information environment relative to a specific question before the final reasoning step occurs.

Local Information Density

Each query occupies a location in semantic space.

Some regions are densely populated with relevant evidence.
Other regions contain little or no institutional information.

Dense Region

- Many nearby chunks
- Small nearest-neighbor distances
- Stable retrieval neighborhoods
- Strong semantic support

Sparse Region

- Few nearby chunks
- Large nearest-neighbor distances
- Unstable retrieval
- Greater risk of unsupported inference

Working Hypothesis

Hallucination risk may increase when a query lies in a region of low local information density.

Evidence Diversity

Ten retrieved chunks do not necessarily represent ten independent pieces of evidence.

The system may return

- ten chunks from the same document,
- multiple copies of the same catalog,
- several nearly identical forms,
- or genuinely diverse evidence from independent sources.

Possible Diversity Measures

- Unique documents
- Unique folders
- Unique document types
- Unique authors or offices
- Unique years

Retrieval Stability

A well-supported question should retrieve similar evidence under small changes in wording.

For example:

- Mechanical Engineering major
- Mechanical Engineering program
- ME degree
- New undergraduate engineering program

High Stability

Small query perturbations retrieve substantially overlapping evidence.

Low Stability

Minor wording changes produce very different retrieval neighborhoods.

Interpretation

Evidence Completeness

A strategic question often requires evidence from multiple domains.

For a proposed academic program, relevant domains may include

- curriculum
- faculty expertise
- enrollment demand
- budget
- facilities
- equipment
- accreditation
- strategic alignment

Coverage Question

Did retrieval find evidence for all required domains, or only a subset?

The Query Observatory

Example Question

What resources would be required to launch a Mechanical Engineering major?

Local information density	9.4 / 10
Evidence diversity	8.8 / 10
Provenance quality	9.1 / 10
Retrieval stability	8.4 / 10
Evidence completeness	7.6 / 10

Expected Reasoning Reliability

High

Inspectable Diagnostics

- Nearest-neighbor distances
- Retrieved-document overlap
- Source-category entropy
- Parser quality
- Provenance completeness
- Domain coverage
- Evidence conflicts

Case Study I: Decision Support

Question:

Should CNU create a Mechanical Engineering major?

ISO does not begin by producing an answer.
Instead it asks:

- What evidence exists?
- Is the evidence sufficient?
- What evidence is missing?
- How reliable is the evidence?

Decision Support

ISO assists human decision makers by assembling, organizing, evaluating, and explaining evidence.
The decision itself remains a human responsibility.

Case Study II: Epistemic Humility

Question:

Should CNU eliminate the French major?

ISO retrieved curriculum documents, catalogs, planning forms, and language placement information.

However, the system explicitly reported missing evidence.

- Enrollment history
- Budget
- Faculty workload
- Strategic planning
- Program review

Correct Conclusion

The available evidence is insufficient to support a recommendation.

A Scientist, Not a Chatbot

Traditional chatbots attempt to answer nearly every question.

ISO attempts something different.

Scientist

- Examine evidence
- State assumptions
- Report uncertainty
- Explain conclusions
- Identify missing data

Desired ISO Behavior

- Retrieve evidence
- Measure evidence quality
- Synthesize findings
- Explain reasoning
- Recommend follow-up

The objective is not artificial certainty.

The objective is trustworthy reasoning.

Toward a Science of Institutional Reasoning

We believe that institutional reasoning can become an observable and measurable engineering discipline.

Future observatory metrics may include

- Local information density
- Evidence diversity
- Evidence completeness
- Provenance quality
- Retrieval stability
- Semantic conflict detection
- Expected reasoning reliability
- Longitudinal corpus evolution

These are properties of the **information environment**, not merely of the language model.

Research Directions

Engineering

- Knowledge objects
- Ontology construction
- Corpus observatories
- Query observatories
- Provenance graphs
- Retrieval benchmarking
- Information quality metrics

Scientific Questions

- Does local information density predict answer quality?
- Does evidence diversity reduce hallucinations?
- How should reasoning reliability be estimated?
- Which observatory metrics best predict trustworthy decisions?
- Can semantic ecosystems be optimized?

The Emerging Architecture

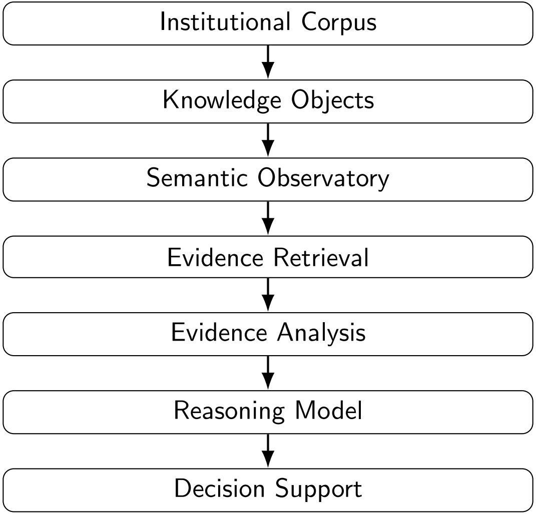

Summary

1. Institutional reasoning begins with evidence.
2. Human intelligence is limited primarily by scale and recall.
3. Retrieval-Augmented Generation is a major advance, but the information environment itself deserves engineering.
4. Knowledge objects store facts.
5. Services derive meaning.
6. Corpus and Query Observatories make reasoning observable, measurable, and explainable.
7. Reliable institutional reasoning is an emergent property of a well-engineered information environment.

Questions?

Discussion

The Institutional Semantic Observatory

Engineering Information Environments for Explainable Institutional AI